



# Using Terraform to define infrastructure as code

Maximilian Hanussek  
High Performance and Cloud Computing Group  
Zentrum für Datenverarbeitung

Applied Bioinformatics Group  
Eberhard Karls University Tübingen

12.10.20  
de.NBI cloud user meeting<sup>3</sup>  
Virtual



- Get a basic understanding of Terraform
- Get an idea how the syntax is working
- Learn what Terraform can be used for
- Create a basic cluster setup (master and compute)



- Introduction
  - Terraform
  - (Virtual-)Cluster
  - Software Stack
- Hands-On
  - Part1: Terraform Basics
  - Part2: Build more complex infrastructures
  - Part3: Basic virtual cluster deployment
  - Part4: VALET



## What is Terraform?

- (Open source) Tool from HashiCorp<sup>1</sup>
- Allows you to maintain infrastructure as Code
- Write configurations as functions and modules
- With all benefits of it (versioning, sharing)
- Provides reproducibility

---

<sup>1</sup><https://www.terraform.io>



## Why Terraform?

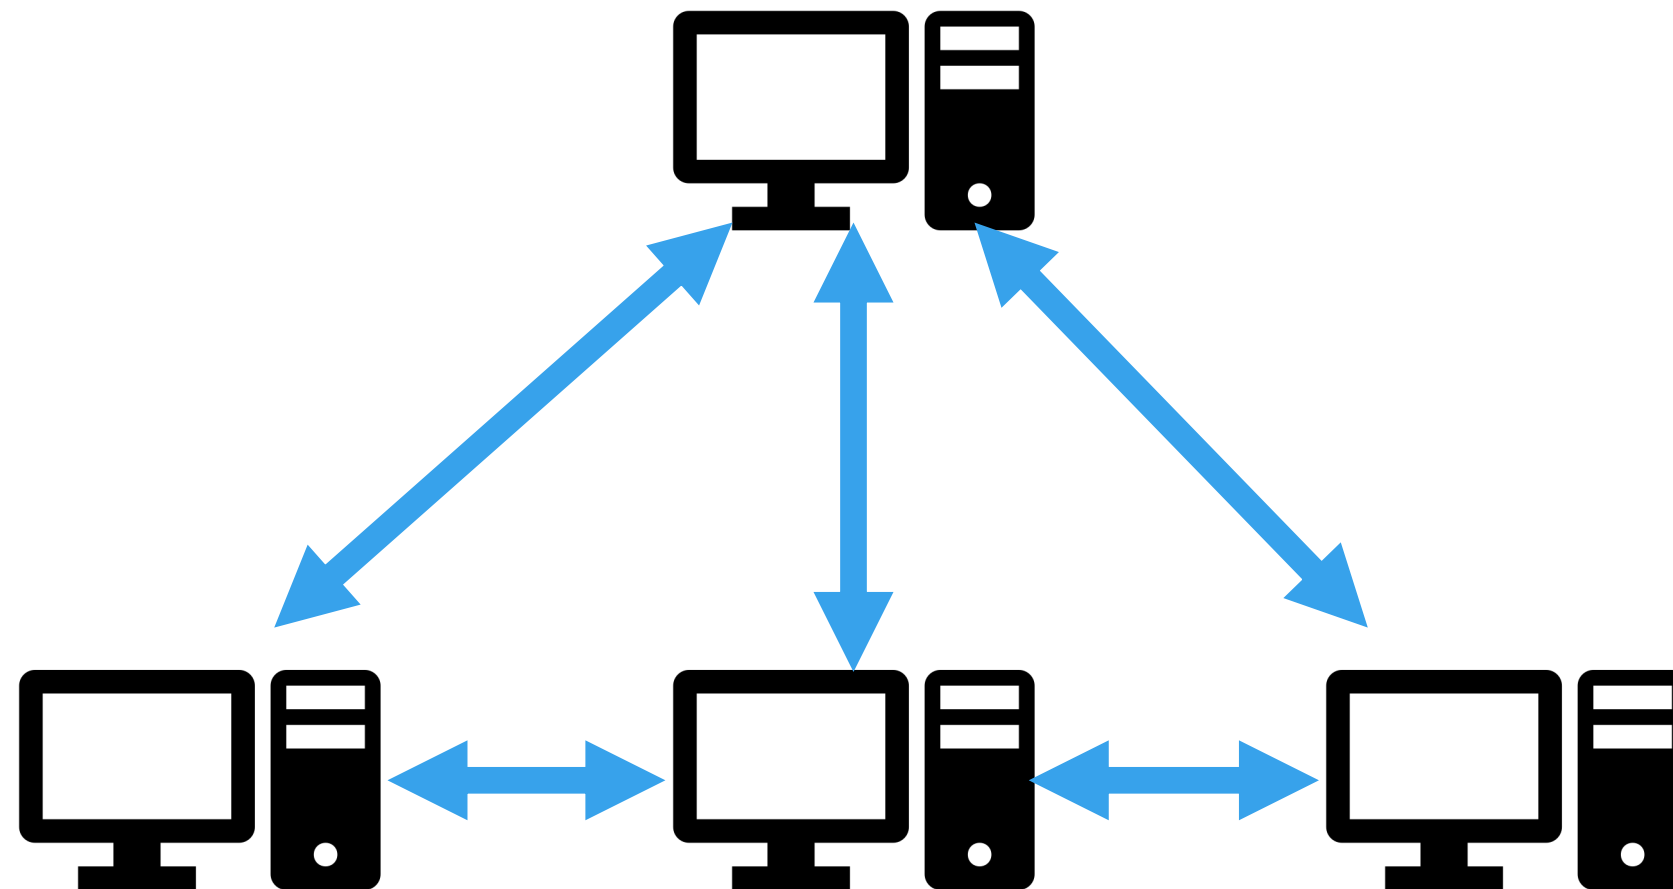
- Easy to install
- Light weighted
- Many tutorials out there
- Comes with a variety of interfaces
- Documentation
- Makes your and our life easier
- Lots of valuable features



A **computer cluster** is a set of loosely or tightly connected computers that work together so that, in many respects, they can be viewed as a single system.<sup>2</sup>

---

<sup>2</sup>[https://en.wikipedia.org/wiki/Computer\\_cluster](https://en.wikipedia.org/wiki/Computer_cluster)

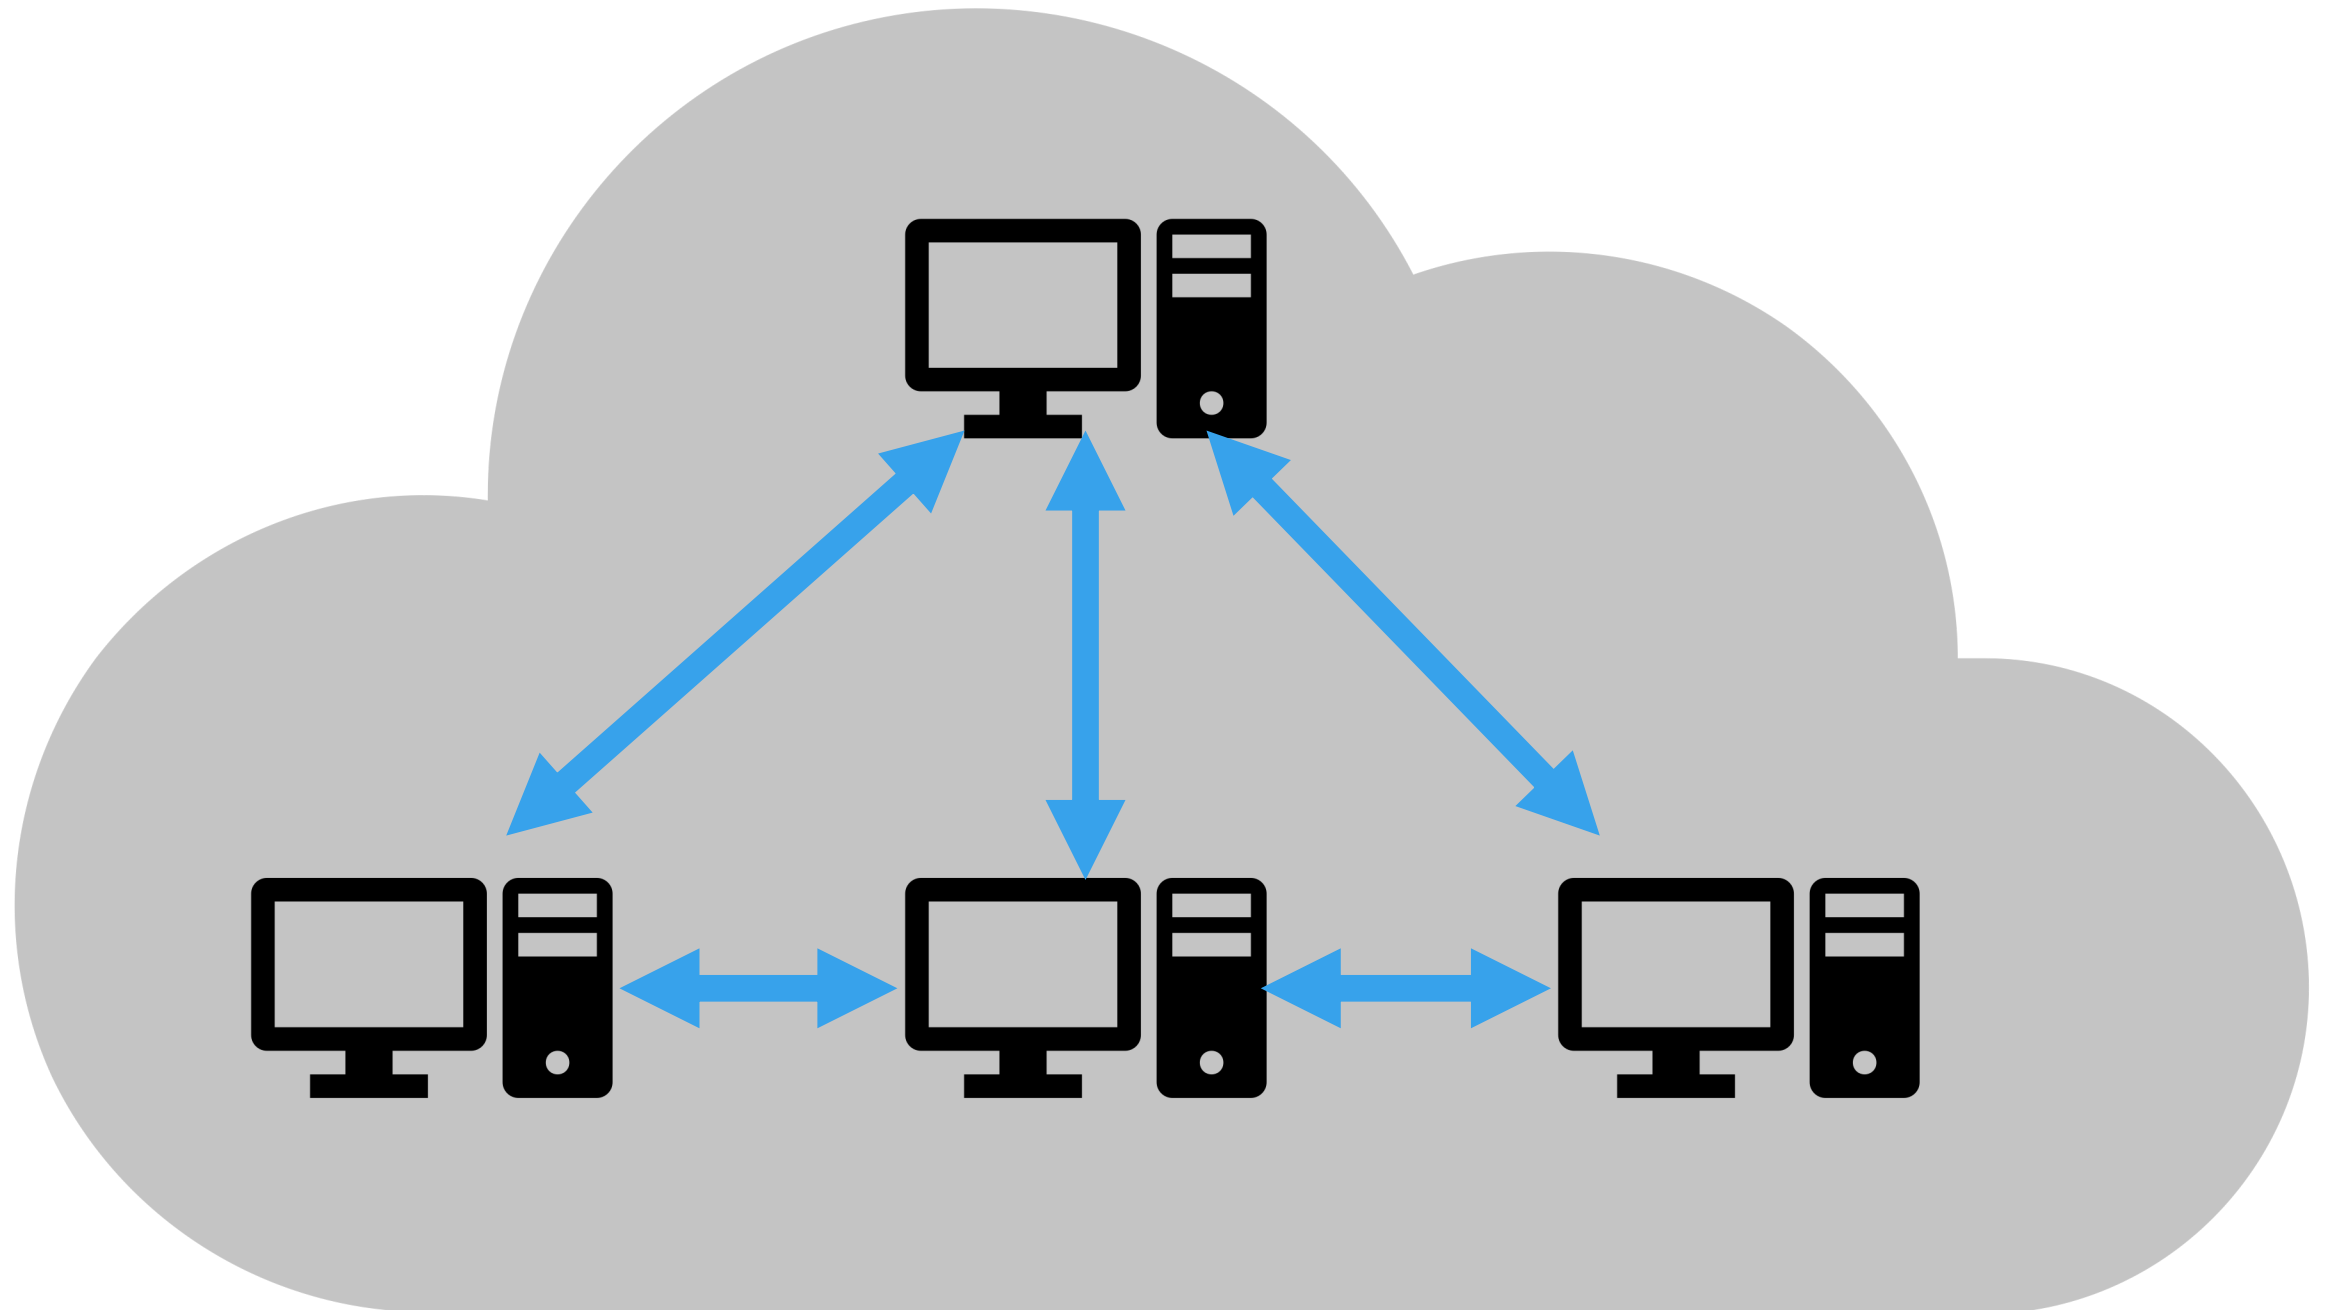




## Use cases:

- Need multiple resources of a single VM
- Your workload can be parallelized
- No need to handle different VMs, just a single cluster
- Run whole pipelines, schedule jobs differently







- Master node(s) (VMs)
- Compute node(s) (VMs)
- Network connection
- (Shared Filesystem)
- (Batch System/Scheduler)
- (Middleware)



## Classical HPC cluster:

- Built up machines
- Configure networks
- Set up shared filesystem
- Configure Scheduler

## Virtual cluster:

- Spawn virtual machines
  - Configure virtual networks
  - Set up shared filesystem
  - Configure Scheduler
- **Use Infrastructure as a Code**



## How to use Terraform?



- Internet connection
- Download and install Terraform (0.13.2)
  - <https://releases.hashicorp.com/terraform/>
- Have a key pair available in the OpenSSH format, if not generate one

```
ssh-keygen -t rsa -b 4096
```


- Download/clone GitHub repo (git clone, wget, curl, via browser)

[https://github.com/MaximilianHanussek/deNBI\\_cloud\\_terraform\\_2020.git](https://github.com/MaximilianHanussek/deNBI_cloud_terraform_2020.git)



- Dashboard Login
  - URL: <https://denbi.uni-tuebingen.de/>



  
**openstack®**

Log in

Keystone Credentials


✓ ELIXIR

If you are not sure which authentication method to use, contact your administrator.

Sign In

<https://denbi.uni-tuebingen.de/>



  
**openstack®**

**Log in**

**Authenticate using**  

Keystone Credentials

If you are not sure which authentication method to use, contact your administrator.

**Domain**  

Default

**User Name**  

workshop\_participant

**Password**

**Region**  

RegionTwo

Sign In





Domain: **Default**

User name: **workshop\_participant**

Password: unicorn42

Region: **RegionTwo**



# Prerequisites– RC File

openstack

Default • deNBI\_user\_meeting • RegionTwo

Project

API Access

Compute

Overview

Instances

Images

Key Pairs

Server Groups

Volumes

Network

Identity

Project / Compute / Overview

## Overview

### Limit Summary

#### Compute

Instances

Used 0 of 1,000

VCPUs

Used 0 of 1,000

RAM

Used 0Bytes of 9.8TB

#### Volume

Volumes

Used 0 of 1,000

Volume Snapshots

Used 0 of 10

Volume Storage

Used 0Bytes of 10TB

#### Network

Floating IPs

Allocated 0 of 50

Security Groups

Used 0 of 10,000

Security Group Rules

Used 0 of 30,000

Networks

Used 0 of 100

Ports

Used 0 of 1,000

Routers

Used 0 of 10

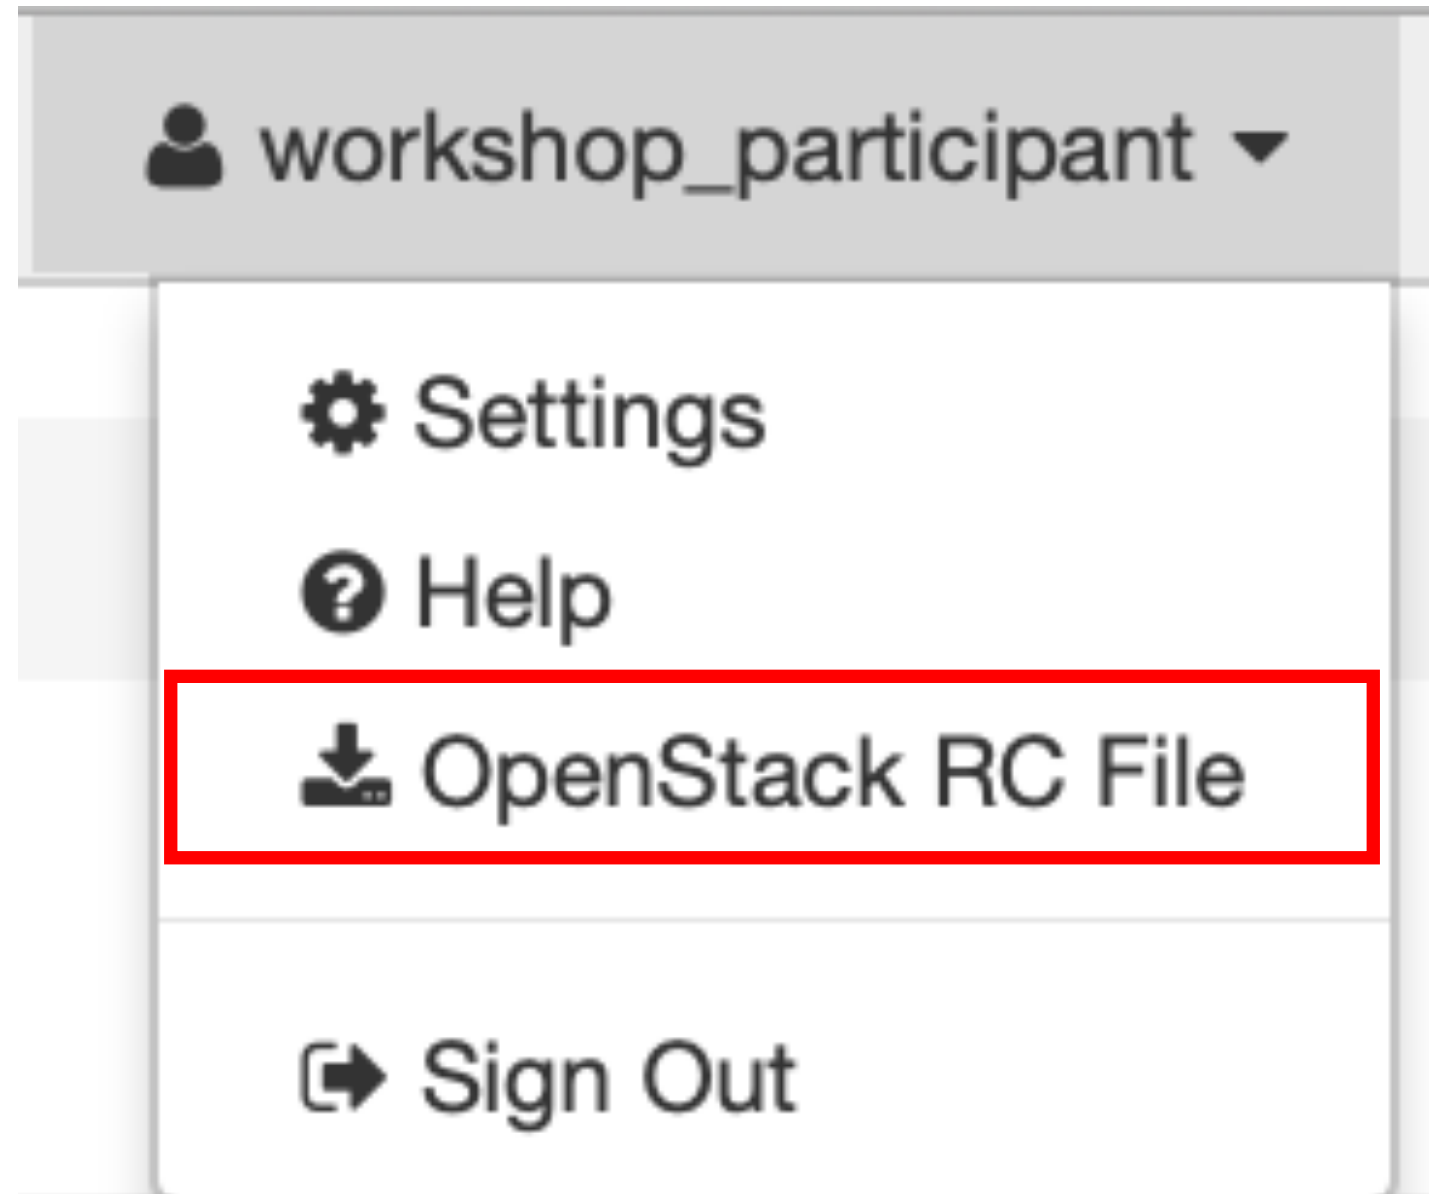
workshop\_participant

Settings

Help

OpenStack RC File

Sign Out





- **Goal:** Understand TF structure and start VM
- Change into:  
`denBI_cloud_terraform_workshop/terraform_workshop_part_1`
- See a bunch of files
  1. `providers.tf`
  2. `key_pair.tf`
  3. `security_group.tf`
  4. `vars.tf`
  5. `main.tf`
  6. `versions.tf`



providers.tf

```
provider "openstack" {}
```



versions.tf

```
terraform {  
  required_providers {  
    openstack = {  
      source = "terraform-providers/openstack"  
    }  
  }  
  required_version = ">= 0.13"  
}
```



## key\_pair.tf

```
resource "openstack_compute_keypair_v2" "workshop_keypair" {  
  name      = var.workshop-key-name  
  public_key = var.public-key  
}
```

- terraform key words
- user chosen name



## security\_group.tf

```
resource "openstack_networking_secgroup_v2"  
"terraform_workshop_sec_group" {
```

```
  name      = var.security-groups[0]  
  description = "Allow SSH access and ping request"  
  delete_default_rules = true
```

```
}
```

- terraform key words
- user chosen name
- key word values





## security\_group.tf (outgoing connections)

```
resource "openstack_networking_secgroup_rule_v2" "egress_public_4" {
  direction      = "egress"
  ethertype      = "IPv4"
  security_group_id =
openstack_networking_secgroup_v2.terraform_workshop_sec_group.id
}
```

```
resource "openstack_networking_secgroup_rule_v2" "egress_public_6" {
  direction      = "egress"
  ethertype      = "IPv6"
  security_group_id =
openstack_networking_secgroup_v2.terraform_workshop_sec_group.id
}
```



## security\_group.tf (SSH)

```
resource "openstack_networking_secgroup_rule_v2"  
"ingress_public_4_ssh" {
```

direction	=	"ingress"
ethertype	=	"IPv4"
protocol	=	"tcp"
port_range_min	=	22
port_range_max	=	22
remote_ip_prefix	=	"0.0.0.0/0"
security_group_id	=	

```
openstack_networking_secgroup_v2.terraform_workshop_sec_group.  
id  
}
```



## security\_group.tf (ICMP)

```
resource "openstack_networking_secgroup_rule_v2"
  "ingress_public_4" {
```

direction	=	"ingress"
ethertype	=	"IPv4"
protocol	=	"icmp"
remote_ip_prefix	=	"0.0.0.0/0"
security_group_id	=	

```
openstack_networking_secgroup_v2.terraform_workshop_sec_group.id
}
```



## vars.tf (volume)

```
variable "cinder-disc-size" {  
  default = 10  
}
```



Needs to be changed



Cloud site specific

```
variable "cinder-storage-backend" {  
  default = "quobyte_hdd"  
}
```

```
variable "volume-name" {  
  default = "maxhanussek-workshop-volume"  
}
```



vars.tf (flavor, name, image)

```
variable "flavors" {  
  type = map  
  default = {  
    "workshop-vm" = "de.NBI default"  
  }  
}
```

```
variable "vm-name" {  
  default = "maxhanussek-workshop-vm"  
}
```

```
variable "workshop-image" {  
  default = "CentOS 7.7 2020-07-07"  
}
```



vars.tf (keys)

```
variable "workshop-key-name" {  
  default = "maxhanussek-keypair"  
}
```

```
variable "public-key" {  
  default = ""  
}
```



vars.tf (security-group, network)

```
variable "security-groups" {  
  default = [  
    "maxhanussek-sec-group"  
  ]  
}
```

```
variable "network" {  
  default = "denbi_uni_tuebingen_external2"  
}
```



## main.tf (image, volume)

```
data "openstack_images_image_v2" "workshop_image" {
  name           = var.workshop-image
  most_recent    = true
}
```

```
resource "openstack_blockstorage_volume_v2" "cinder_volume" {
  name           = var.volume-name
  size           = var.cinder-disc-size
  volume_type    = var.cinder-storage-backend
}
```





## main.tf (instance)

```
resource "openstack_compute_instance_v2" "workshop_vm" {
```

```
    name = var.vm-name
    flavor_name = var.flavors["workshop-vm"]
    image_id = data.openstack_images_image_v2.workshop_image.id
    key_pair = openstack_compute_keypair_v2.workshop_keypair.name
    security_groups = var.security-groups
```

```
    network {
        name = var.network
    }
```



## main.tf (instance)

```
block_device {
  uuid =
data.openstack_images_image_v2.workshop_image.id
  source_type = "image"
  destination_type = "local"
  boot_index = 0
  delete_on_termination = true
}
```

```
block_device {
  uuid =
openstack_blockstorage_volume_v2.cinder_volume.id
  source_type = "volume"
  destination_type = "volume"
  boot_index = -1
  delete_on_termination = true
}
```

}



## Use Terraform

- Initialize Terraform in Terraform directory  
(deNBI\_cloud\_terraform\_2020/terraform\_workshop\_part\_1)

```
terraform init
```



**Initializing the backend...**

**Initializing provider plugins...**

...

\* provider.openstack: version = "~> 1.31.0"

**Terraform has been successfully initialized!**

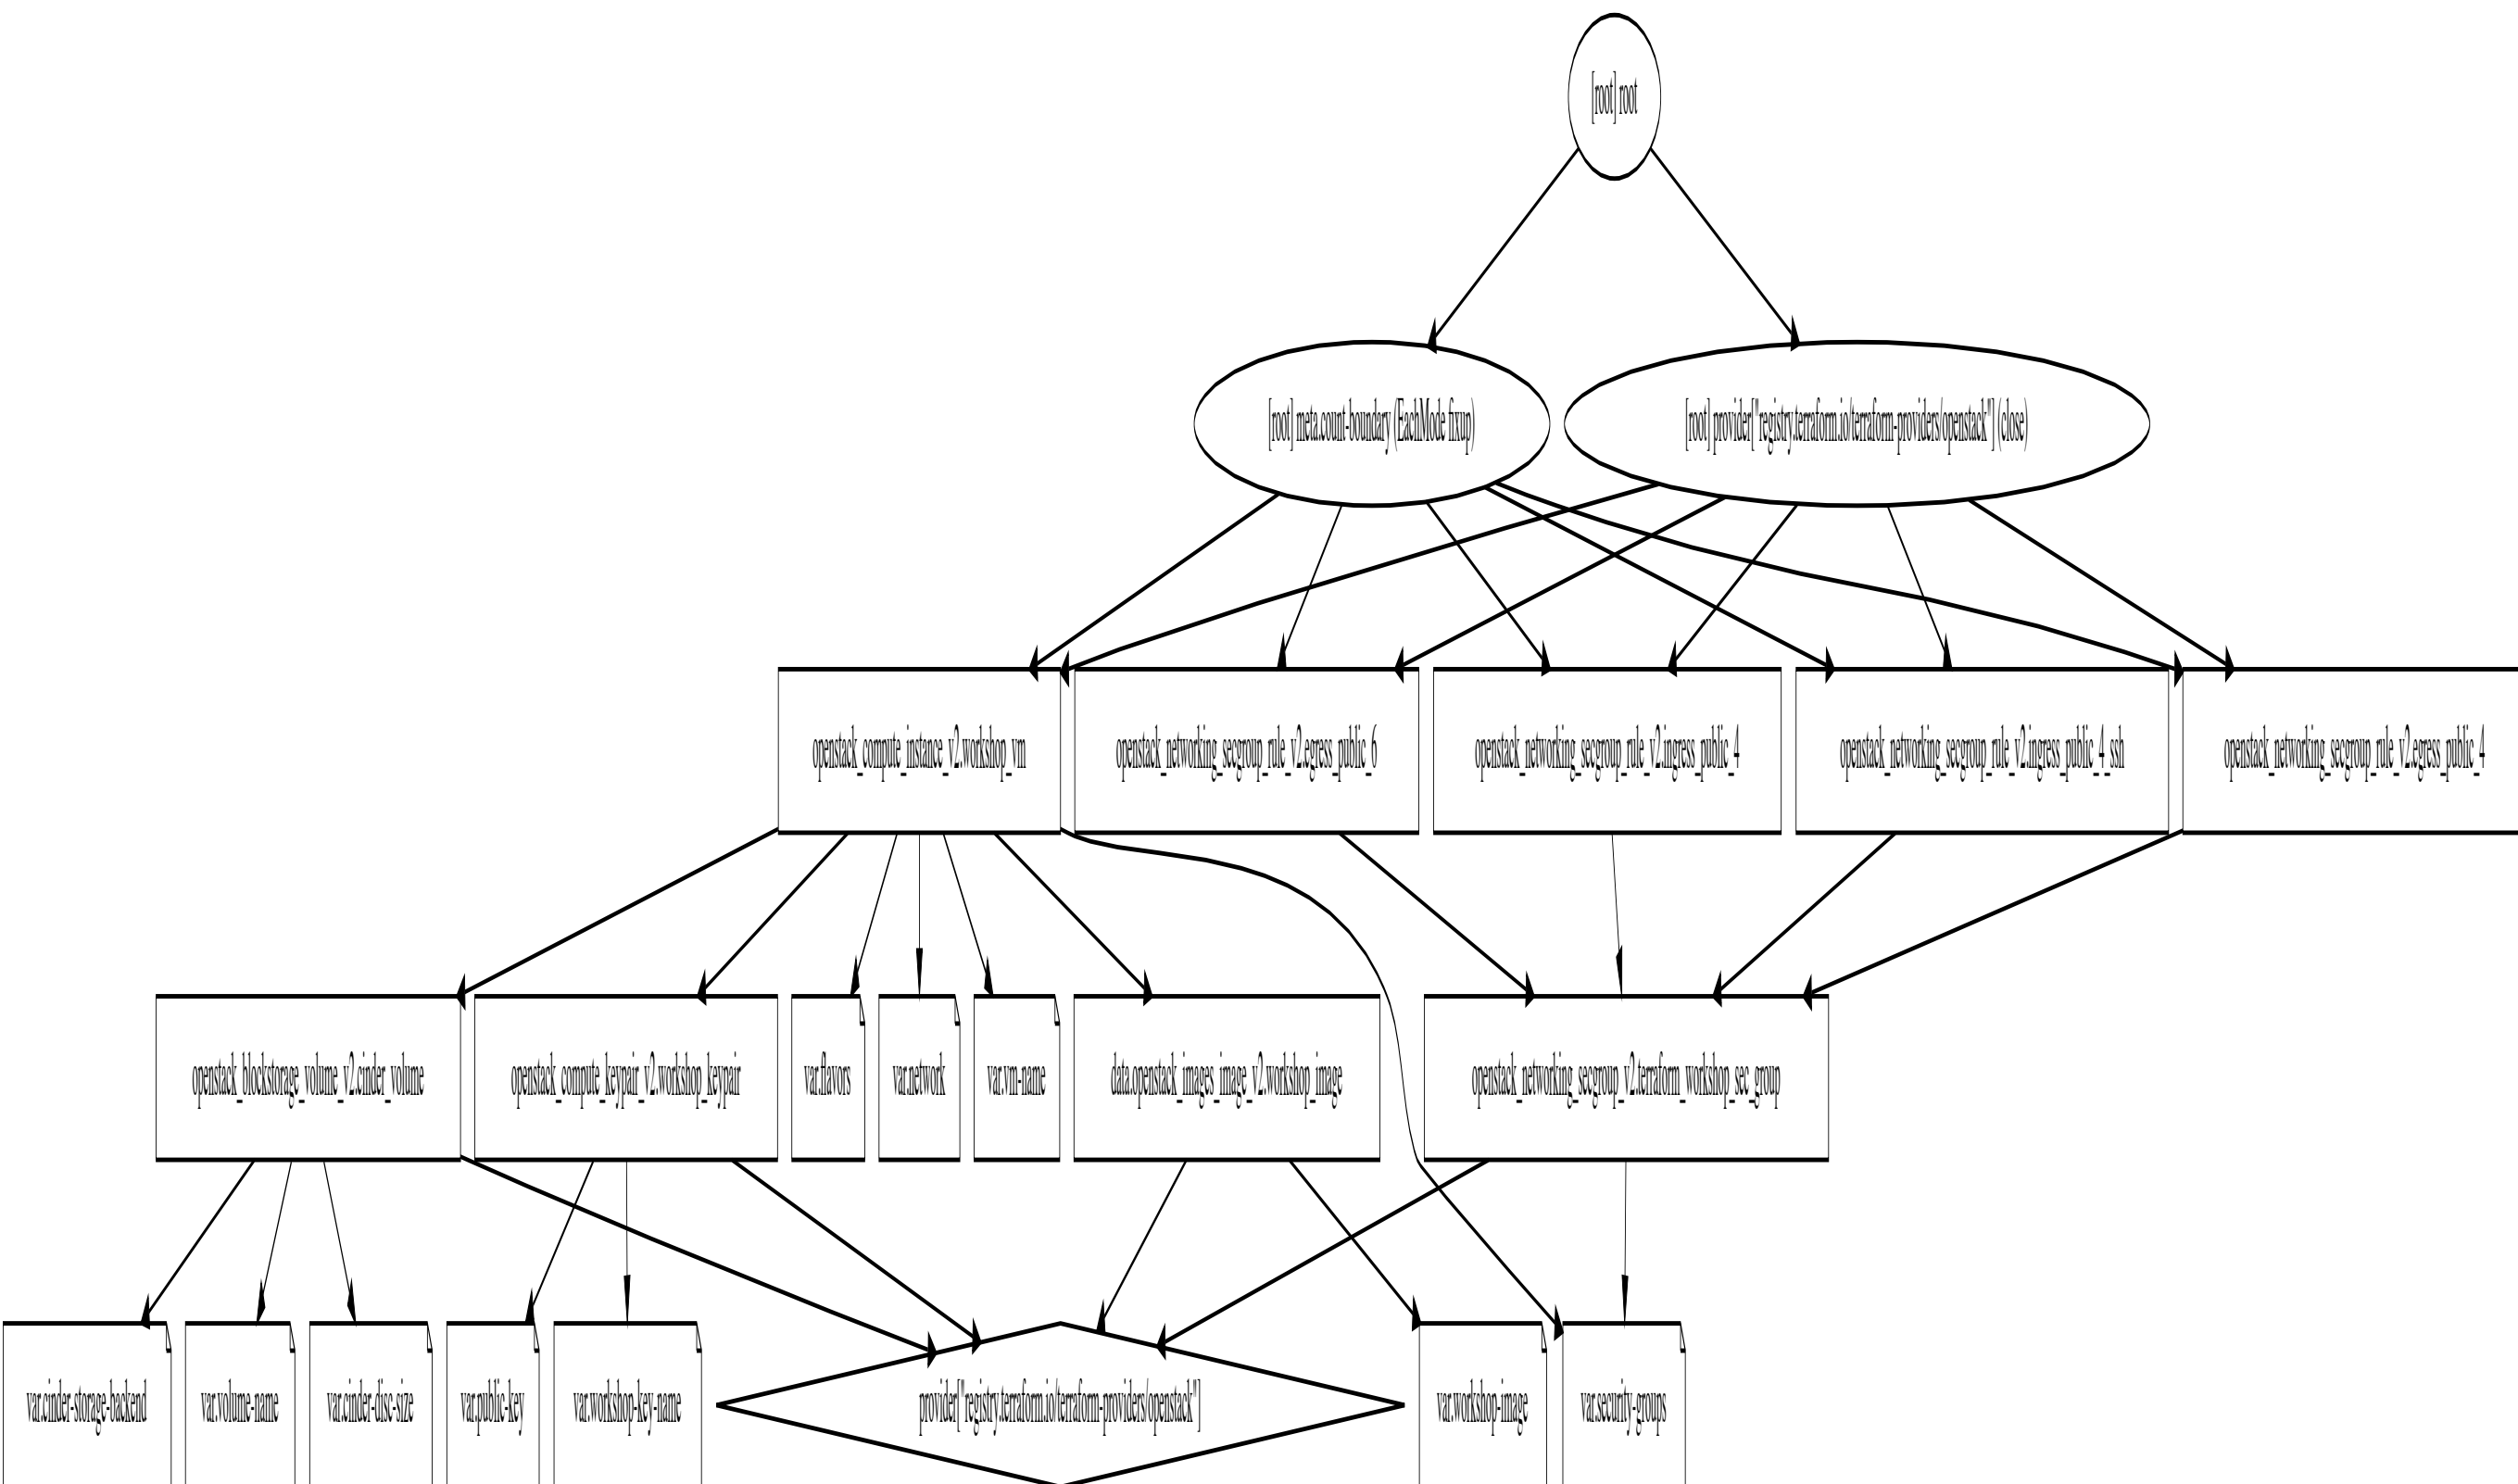
...



## Use Terraform

- Run Terraform graph  
(denBI\_cloud\_terraform\_2020/terraform\_workshop\_part\_1)

```
terraform graph | dot -Tsvg > graph.svg
```





## Use Terraform

- Run Terraform plan  
(denBI\_cloud\_terraform\_2020/terraform\_workshop\_part\_1)

```
terraform plan
```



**Error: One of 'auth\_url' or 'cloud' must be specified**

```
on providers.tf line 1, in provider "openstack":  
  1: provider "openstack" {}
```





## Use Terraform

- Source OpenStack credentials  
(deNBI\_cloud\_terraform\_2020/terraform\_workshop\_part\_1)

```
source deNBI_user_meeting-openrc.sh
```

Password:

```
terraform plan
```



```
# openstack_blockstorage_volume_v2.cinder_volume will be
created
+ resource "openstack_blockstorage_volume_v2"
"cinder_volume" {
    + attachment          = (known after apply)
    + availability_zone    = (known after apply)
    + id                  = (known after apply)
    + metadata             = (known after apply)
    + name                 = "maxhanussek-workshop-volume"
    + region               = (known after apply)
    + size                 = 10
    + volume_type          = "quobyte_hdd"
}
```



**Plan: 8 to add, 0 to change, 0 to destroy.**



## Use Terraform

- Run Terraform apply (confirm with 'yes')  
(denBI\_cloud\_terraform\_2020/terraform\_workshop\_part\_1)

```
terraform apply
```

Apply complete! Resources: 8 added, 0 changed, 0 destroyed.

Check the Dashboard for the created resources

<https://denbi.uni-tuebingen.de>



## Use Terraform

- Run Terraform destroy (confirm with 'yes')  
(denBI\_cloud\_terraform\_2020/terraform\_workshop\_part\_1)

```
terraform destroy
```

Destroy complete! Resources: 8 destroyed.

Check the Dashboard for the destroyed resources

<https://denbi.uni-tuebingen.de>



### Use Terraform

- Change to: `terraform_workshop_part_2`
- Initialize Terraform in Terraform directory  
(`deNBI_cloud_terraform_2020/terraform_workshop_part_2`)

```
terraform init
```



### Use Terraform

- Goal: Start multiple (similar) VMs
- Add new variable to end of vars.tf

```
variable "node-count" {  
    default = 3  
}
```



## Use Terraform

- Change parts of main.tf (Volume resource)

```
resource "openstack_blockstorage_volume_v2"
"cinder_volume" {
  count = var.node-count
  name = "${var.volume-name}-${count.index}"
  size = var.cinder-disc-size
  volume_type = var.cinder-storage-backend
}
```





### Use Terraform

- Change parts of main.tf ( Instance resource)

```
resource "openstack_compute_instance_v2" "workshop_vm" {  
  count      = var.node-count  
  name       = "${var.vm-name}-${count.index}"  
  flavor_name = var.flavors["workshop-vm"]  
  image_id    = data.openstack_images_image_v2.workshop_image.id  
  key_pair    = openstack_compute_keypair_v2.workshop_keypair.name  
  security_groups = var.security-groups  
  
  network {  
    name = var.network  
  }  
}
```



### Use Terraform

- Change parts of main.tf ( Block device)

```
block_device {  
  uuid =  
element(openstack_blockstorage_volume_v2.cinder_volume.*.id  
, count.index)  
  source_type          = "volume"  
  destination_type     = "volume"  
  boot_index           = -1  
  delete_on_termination = true  
}  
}
```



### Use Terraform

- Change to: `terraform_workshop_part_2`
- Run `terraform plan` & `terraform apply`  
(`deNBI_cloud_terraform_2020/terraform_workshop_part_2`)

`terraform plan`

`terraform apply`



### Use Terraform

- Changed network variable in vars.tf

```
variable "network" {  
    default = "denbi_uni_tuebingen_internal"  
}
```



### terraform.tfstate

- Have a look at terraform.tfstate file
  - Holds all information about the infrastructure
  - !!! Holds also your credentials
  - !!! Holds also keys
- Be careful if you share these file (Git, ...)

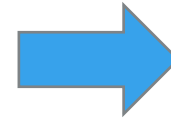




## Use Terraform

- Change the infrastructure
- Open `vars.tf` file

```
variable "node-count" {  
    default = 3  
}
```



```
variable "node-count" {  
    default = 4  
}
```



`terraform plan`

`Plan: 2 to add, 0 to change, 0 to destroy.`

`terraform apply`

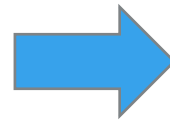
`Apply complete! Resources: 2 added, 0 changed, 0 destroyed.`



### Use Terraform

- Change the infrastructure again
- Open `vars.tf` file

```
variable "node-count" {  
  default = 4  
}
```



```
variable "node-count" {  
  default = 2  
}
```





`terraform plan`

`Plan: 0 to add, 0 to change, 4 to destroy.`

`terraform apply`

`Apply complete! Resources: 0 added, 0 changed, 4 destroyed.`



### Use Terraform

- Change the infrastructure by accident
- Go to the dashboard and delete a VM
- (Please take one of your own ones!)

terraform plan

terraform apply



`terraform plan`

`Plan: 2 to add, 0 to change, 0 to destroy.`

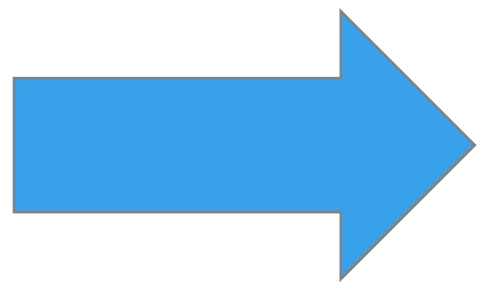
`terraform apply`

`Apply complete! Resources: 2 added, 0 changed, 0 destroyed.`



### Use Terraform

- Changing the infrastructure through
- Variable changes (remove or add)
- Handling of deletions from outside



Terraform makes it easy to recover from unwanted changes or scaling up and down the infrastructure



## Use Terraform

- Run terraform destroy

**terraform destroy**

- Check terraform.tfstate again

```
{  
  "version": 4,  
  "terraform_version": "0.13.2",  
  "serial": 183,  
  "lineage": "8c08c491-29cd-9f7a-b79c-5f607a59374d",  
  "outputs": {},  
  "resources": []  
}
```



### Use Terraform

- Change to: `terraform_workshop_part_3`
- Initialize Terraform in Terraform directory  
(`deNBI_cloud_terraform_2020/terraform_workshop_part_3`)

```
terraform init
```



## Use Terraform

- Goal: Build a cluster structure
- Cluster setup:
  - 3 VMs (1 Master node, 2 Compute nodes)
  - 3 Volumes
  - 2 different networks
  - 2 different flavors
  - 2 different images



## vars.tf

```
variable "volume-name" {
  type = map
  default = {
    "master" = "maxhanussek-workshop-volume-master"
    "compute" = "maxhanussek-workshop-volume-compute"
  }
}
```

```
variable "vm-name" {
  type = map
  default = {
    "master" = "maxhanussek-workshop-vm-master"
    "compute" = "maxhanussek-workshop-vm-compute"
  }
}
```





## vars.tf

```
variable "flavors" {
  type = map
  default = {
    "master" = "de.NBI small"
    "compute" = "de.NBI default"
  }
}
```

```
variable "workshop-image" {
  type = map
  default = {
    "master" = "CentOS 7.7 2020-07-07"
    "compute" = "CentOS 7.8 2020-07-07"
  }
}
```



## vars.tf

```
variable "network" {
  type = map
  default = {
    "master" = "denbi_uni_tuebingen_external2"
    "compute" = "denbi_uni_tuebingen_internal"
  }
}
```

```
variable "private-key-path" {
  default = "/path/to/private/key"
}
```

```
variable "internal-key-name" {
  default = "maxhanussek-internal-key"
}
```



## key\_pair.tf

```
resource "tls_private_key" "internal_connection_key" {
  algorithm = "RSA"
  rsa_bits  = 4096
}

resource "openstack_compute_keypair_v2"
"internal_connection_key_openstack" {
  name      = var.internal-key-name
  public_key =
tls_private_key.internal_connection_key.public_key_openssh
}
```



## main.tf

```
data "openstack_images_image_v2" "workshop_image_master" {
  name           = var.workshop-image["master"]
  most_recent    = true
}
```

```
data "openstack_images_image_v2" "workshop_image_compute" {
  name           = var.workshop-image["compute"]
  most_recent    = true
}
```



## main.tf

```
resource "openstack_compute_instance_v2" "workshop_vm_master" {
  name          = var.vm-name["master"]
  flavor_name   = var.flavors["master"]
  image_id      =
data.openstack_images_image_v2.workshop_image_master.id
  key_pair      =
openstack_compute_keypair_v2.workshop_keypair.name
  security_groups = var.security-groups

  network {
    name = var.network["compute"]
  }

  network {
    name          = var.network["master"]
    access_network = true
  }
}
```



## main.tf (provisioner)

```

provisioner "file" {
  content = tls_private_key.internal_connection_key.private_key_pem
  destination = "~/.ssh/connection_key.pem"

  connection {
    type          = "ssh"
    private_key   = file(var.private-key-path)
    user          = "centos"
    timeout       = "5m"
    host          = self.access_ip_v4
  }
}

```



## main.tf (provisioner)

```
provisioner "file" {  
  source      = "hello_world.txt"  
  destination = "/home/centos/hello_world.txt"
```

```
  connection {  
    type          = "ssh"  
    private_key   = file(var.private-key-path)  
    user          = "centos"  
    timeout       = "5m"  
    host          = self.access_ip_v4  
  }  
}
```



## main.tf (provisioner)

```
provisioner "remote-exec" {  
  script = "set_internal_private_key_permissions.sh"
```

```
  connection {  
    type          = "ssh"  
    private_key   = file(var.private-key-path)  
    user          = "centos"  
    timeout       = "5m"  
    host          = self.access_ip_v4  
  }  
}
```





## main.tf

```
provisioner "remote-exec" {
  script = "mount_cinder_volumes.sh"

  connection {
    type          = "ssh"
    private_key   = file(var.private-key-path)
    user          = "centos"
    timeout       = "5m"
    host          = self.access_ip_v4
  }
}
```



## Use Terraform

- Change to: `terraform_workshop_part_3`
- Run `terraform plan` and `terraform apply`  
(`deNBI_cloud_terraform_2020/terraform_workshop_part_3`)

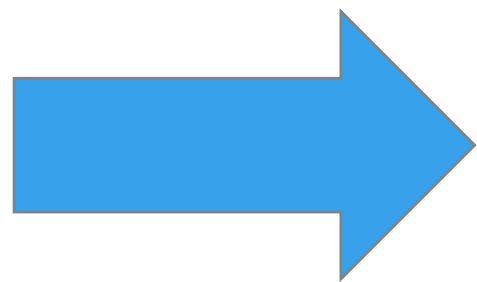
`terraform plan`

`terraform apply`

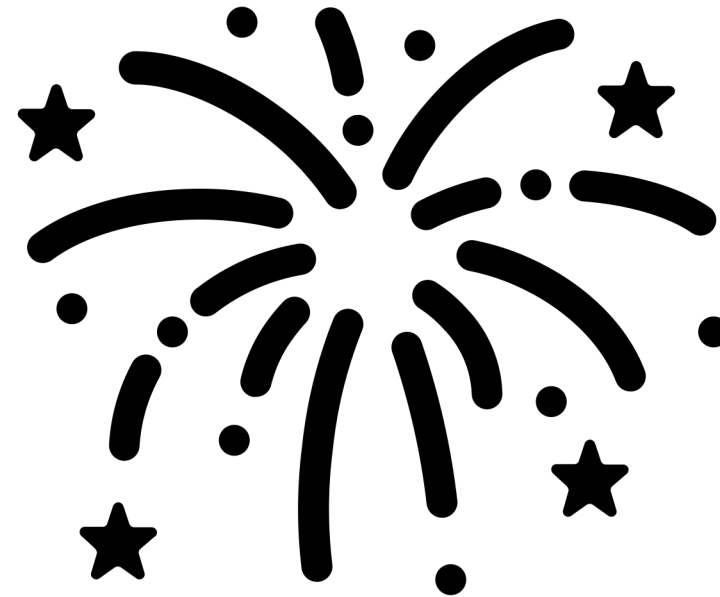


### Use Terraform

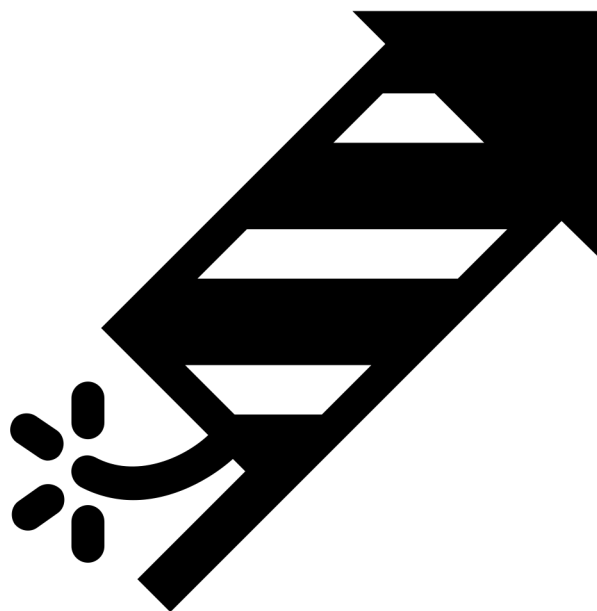
- Make use of the map module
- Let TF create “one time” keys for you
- Provisioner module can be used for post mods



Terraform makes it more convenient to deploy more complex infrastructures in an automated way and also be able to handle it.



Congratulations we have deployed  
a virtual cluster infrastructure





### Use Terraform

- Goal: Build an advanced virtual cluster (VALET)
- Features:
  - Batch System (PBS TORQUE)
  - Middleware UNICORE (incl. Workflow engine)
  - Monitoring system (Zabbix)
  - Add and remove nodes without downtime
  - Automated load based scaling available

<https://github.com/MaximilianHanussek/VALET>